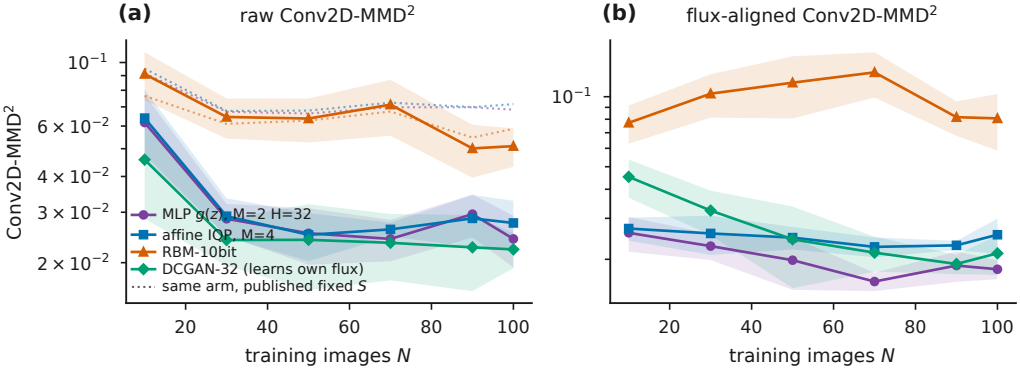

Mean ± 1 sd over 5 seeds; MLP $g(z)$ n=5, affine IQP n=5, RBM-10bit n=5 at N=100. Solid = resampled S, dotted = the published fixed S (panel (a) only: the aligned metric in (b) is identical under either convention by construction). DCGAN has one curve because it learns a per-image flux and has no fixed S to swap; the same 200 held-out frames and the same v11 metrics are used throughout.